

- DropKast — Budget & Operations Plan
 - TL;DR — The numbers you'll be asked
 - Phase 1 — Build & launch (\$0 → first 50 artists)
 - One-time costs
 - Monthly recurring (soft launch tier)
 - Engineering effort
 - Phase 2 — Real product (50 → 500 artists)
 - Monthly recurring
 - Revenue model assumptions
 - Phase 3 — Growth (500 → 5,000 artists)
 - When to hire
 - Phase 4 — Scaled (5K → 50K artists, DistroKid territory)
 - What scaled product looks like
 - AI maintenance costs — projected
 - Per-user cost breakdown at typical use
 - Cost guardrail (already built)
 - Switching to free providers
 - Image / video at scale
 - Security & compliance budget
 - Phase 2 (50–500 artists)
 - Phase 3+ (real B2B contracts)
 - Server / cloud plan recommendations
 - Recommended stack (fits the architecture)
 - Don't do these (common mistakes)
 - "How much for AI to maintain most of this site?"
 - What "AI maintaining the site" means concretely
 - Cost projection for AI-assisted maintenance
 - The honest bound
 - Recommended decision framework
 - Sources

DropKast — Budget & Operations Plan

How much it costs to **build, launch, and run** DropKast at various scales. All numbers are USD, current as of April 2026, and assume you're following the architecture in this repo (Vercel + Supabase + Anthropic + Stripe + RouteNote).

Audience: non-technical leadership making a build/buy/scale decision. If you want the engineering breakdown of WHAT to integrate, see [CONNECTIONS_NEEDED.md](#).

TL;DR — The numbers you'll be asked

Scenario	Monthly cost	Per active artist	Notes
Soft launch (50 artists, demo mode)	~\$80/mo	\$1.60	Free Vercel/Supabase tiers + small Anthropic spend
Real product (500 artists, real distro)	~\$1,200/mo	\$2.40	Supabase Pro, Vercel Pro, Stripe Connect, RouteNote partner
Growth (5,000 artists)	~\$8,500/mo	\$1.70	Supabase Team, Vercel Pro at scale, more AI spend
Scaled (50,000 artists, à la DistroKid)	~\$60,000/mo	\$1.20	Enterprise plans, dedicated DSP infra, SOC2, support team

The per-artist cost goes **down** with scale because most fixed costs (database, distribution agreements, support tooling) don't grow linearly with users. Variable costs (AI calls, payouts, audio storage) do grow but stay predictable.

Phase 1 — Build & launch (\$0 → first 50 artists)

One-time costs

Item	Cost	Notes
RIAA ISRC registrant	\$95	One-time. Mints unlimited ISRCs forever.
GS1 UPC company prefix	~250+ 50/yr	Annual renewal scales with company revenue.
Domain (e.g. dropkast.com)	\$12/yr	Cloudflare Registrar at cost.
Legal — Terms, Privacy, DPA	1, 500–5,000	Use a SaaS-tailored lawyer (Latham, etc.)
Trademark filing	\$250	DIY USPTO. \$1,500 with attorney.
Total one-time	~\$2,000	Bare minimum to launch legitimately.

Monthly recurring (soft launch tier)

Item	Cost	Notes
Vercel Pro	\$20/seat/mo	1 seat at launch.
Supabase Free	\$0	Up to 500MB DB, 50K MAU — fine for first ~50 artists.
Anthropic API	\$20–50/mo	Assuming ~50 active users with ~10 AI sessions each.
Cloudflare DNS + WAF	\$0	Free tier covers DDoS + basic WAF.
Email (Resend or Postmark)	0–15/mo	Free up to 3K emails.
Monthly total	~\$40–80/mo	Cheap.

Engineering effort

Phase	Time	Status
Phase 0 (auth + DB + Zod)	5 days	✅ Code ready, need Supabase project
Phase 1 (DSP delivery + ISRC + state machine)	2 weeks	⚠️ Code ready except RouteNote integration
Phase 2 (AI everywhere)	1 week	✅ Done
Phase 3 (Stripe Connect splits)	1 week	🟡 Schema ready, needs onboarding flow
Phase 5 (polish layer)	5 days	✅ Done

Real elapsed time assuming external blockers: 6–8 weeks (RouteNote partnership is the long pole).

Phase 2 — Real product (50 → 500 artists)

You're charging money. SLAs matter. Splits actually pay out. DSP delivery is real.

Monthly recurring

Item	Cost	Why
Vercel Pro	\$20/mo	Bandwidth + functions still under 1TB / 1M invocations.
Supabase Pro	\$25/mo	8GB DB + 100GB storage + 100K MAU.
Anthropic API	\$200–500/mo	500 users × ~\$0.40–1 of AI use/mo. Prompt caching keeps it down.
Cloudflare R2 (audio storage)	\$50/mo	~3TB stored at \$0.015/GB-month. Zero egress is the win.
RouteNote partner share	0–10% of distributed royalties	They take a cut, not a flat fee.
Stripe Connect	\$2/active payee/mo + 2.9%+30¢ per txn	If 250 active payees: \$500/mo platform fee.
Sentry Team	\$26/mo	50K errors/mo.
PostHog (free tier)	\$0	First 1M events free.
Resend (transactional email)	\$20/mo	50K emails.
Domain + Cloudflare	\$1/mo amortized	Negligible.
Monthly subtotal	~\$850–1,100	Plus payment processing on transfers.

Item	Cost	Why
+ Stripe processing fees	~3% of GMV	If 50K royalties moved through : 1,500.
Monthly total	~\$1,000–2,500	Most variability is Stripe fees.

Revenue model assumptions

Three ways DropKast can make money — in order of how repeatable they are:

1. **Premium subscription** — $15/\text{mo per artist} \times 500 \text{ artists} \times 50/15 = \$3,750/\text{mo MRR}$. Covers the entire monthly stack.
2. **Per-release fee** — $9.99/\text{release like RouteNote Premium} \times 500 \text{ artists} \times 1.5 \text{ releases/yr avg} \times 9.99 = \$7,500/\text{yr}$.
3. **Royalty share** — 10% of artist royalties (RouteNote-style). Best long-term, hardest to pitch. If platform GMV is 200K/yr : 20K/yr revenue.

Hybrid pricing (recommended): free tier with 10% royalty share + \$15/mo premium that drops the share to 0%. Mirrors RouteNote → DistroKid migration path.

Phase 3 — Growth (500 → 5,000 artists)

Item	Cost	Why
Vercel Pro	\$40–80/mo	2–4 seats; usage may push you toward Enterprise convo.
Supabase Team	\$599/mo base	Higher caps, daily backups, point-in-time recovery.
Anthropic API	\$1,500–4,000/mo	5K users × ~\$0.30–0.80 AI use. Heavy use of Haiku for cheap tasks.
Free chat providers (NVIDIA / Groq / Cerebras)	\$0	Free tier covers fallback.
Cloudflare R2	\$200/mo	~13TB stored.
Stripe Connect platform fees	\$2K/mo	~1K active payees.
Sentry Business	\$80/mo	More errors, more retention.
PostHog (paid)	\$200/mo	~600M events at \$0.00031/event.
Customer support tool (Intercom / Crisp)	\$200/mo	You'll need this.
Monthly subtotal	~\$5K	
Stripe processing	~\$3.5K/mo	At \$120K GMV.
Monthly total	~\$8.5K	

When to hire

- **First eng hire** at ~500 paying artists. Senior full-stack, \$150K–200K base + equity.
- **First support hire** at ~1,000 paying artists. Part-time → full-time at 5K. ~\$60K base.
- **Eng #2 (backend / infra)** at ~5K artists. Same range as #1.
- **A&R coordinator** (the human side of A&R) at ~2K artists. ~\$80K.

Phase 4 — Scaled (5K → 50K artists, DistroKid territory)

This is where you stop being a startup and start being an infrastructure company.

Item	Cost	Why
Vercel Enterprise	\$5K–15K/mo	Custom SLA, higher concurrency caps, dedicated support.
Supabase Team / Enterprise	\$99–2K/mo	Or self-hosted Postgres on AWS RDS with Read Replicas (~\$1K/mo).
Anthropic API	\$10K–25K/mo	At scale, negotiate volume contracts with Anthropic.
Cloudflare Stream + R2	\$1K–3K/mo	Audio CDN delivery scales linearly with plays.
Stripe Connect	\$20K–50K/mo	Per-account fees + processing fees on volume.
SOC2 audit (first time)	20K–50K(one – time)+ 15K/yr	Required for B2B contracts.
Pen test	\$15K/yr	At minimum annually, more often pre-launch.
Engineering team (5–8 people)	\$1M–1.5M/yr in salary	Senior eng is \$200K base avg.
Support team (3–5 people)	\$250K/yr	At \$50–80K each.
Monthly platform total	~\$45K–80K/mo	Plus salaries.

What scaled product looks like

- **Self-hosted Postgres** with read replicas in 2 regions. ~\$2K/mo on AWS RDS.
- **Dedicated DSP delivery infra** — own DDEX pipeline, no longer dependent on RouteNote/SonoSuite. ~\$5K/mo infra + 6 months engineering.
- **24/7 on-call rotation** for the eng team.
- **SOC2 Type II + ISO 27001** for B2B contracts.

AI maintenance costs — projected

The user asked specifically about AI cost at scale. This is the most variable line item. Worth understanding deeply.

Per-user cost breakdown at typical use

A typical artist's AI usage in a month:

Action	Frequency	Model	Tokens (cached)	Cost
5 AI assistant chat sessions @ 10 messages each	50 messages	Sonnet	~80K input (cached) + 30K output	\$0.74
2 A&R critiques	2	Sonnet	~5K input + 2K output	\$0.06
1 campaign strategy generation	1	Sonnet	~3K input + 1K output	\$0.03
5 viral idea generations	5	Haiku	~5K input + 4K output	\$0.01
10 image generations (Flux Schnell)	10	Flux	n/a	\$0.03
Per-artist monthly AI cost				~\$0.87

With prompt caching the input cost is dramatically lower (90% discount on cached reads). Without caching: ~\$2.50/artist/month.

Cost guardrail (already built)

DropKast enforces a $1/dayper-user * AIbudgetbydefault('AI_{DAILY_BUDGET_CENTS} = 100')$. Anyone exceeding gets a 429 with a friendly "budget reset at midnight UTC" message. This caps your worst case exposure to $\$30/artist/month$ even if a user goes wild.

Switching to free providers

If your premium tier doesn't take off, route artists to free providers (NVIDIA/Groq/Cerebras) by default. This drops AI cost from \$0.87/artist to **near-zero** at the cost of:

- No tool use (assistant can't read your data; it's just a chatbot).
- Lower-quality A&R critiques (Llama 3.3 < Claude Sonnet).
- Free tier rate limits sometimes hit (40 RPM on NVIDIA, etc.).

Image / video at scale

Today these aren't wired. Once they are:

- 10 covers $\times 0.003(FluxSchnell) = 0.03/artist$
- 1 video $\times 0.029/sec \times 30sec = 0.87$ for one Kling music video

If artists generate 10 videos/month each at this rate: 8.70/*artist/month*. *
*This is why image/video is the natural Premium tier feature * *—it gates the highest —
cost ops behind the 15/month subscription.

Security & compliance budget

Don't skip these. One leaked database is a regulatory nightmare.

Phase 2 (50–500 artists)

Item	Cost	Frequency
HTTPS / TLS	\$0	Vercel handles.
Encryption at rest	\$0	Supabase default.
Backup & PITR	included in Supabase Pro	\$25/mo.
Cloudflare WAF + DDoS	0free/250 Business	Annual.
2FA on admin accounts	\$0	Use Authy / 1Password.
Phase 2 total	~\$0–250/mo	

Phase 3+ (real B2B contracts)

Item	Cost	Frequency
SOC2 Type II audit	\$20K–50K	Year 1. \$15K/yr ongoing.
Vanta or Drata (compliance tooling)	\$7K–15K/yr	Required for SOC2 prep.
Penetration test	\$10K–20K	Annually.
Bug bounty (HackerOne)	\$5K–20K/yr	When ready.
Cyber insurance	\$5K–15K/yr	\$1M coverage.
Phase 3+ total	~\$50K/yr	Plus per-incident if breach.

Server / cloud plan recommendations

Recommended stack (fits the architecture)

Service	What it does	Why this one
Vercel Pro	Hosts the React app + serverless API	Already configured. Deploys on git push.

Service	What it does	Why this one
Supabase Pro	Postgres + Auth + Storage	Postgres native (Drizzle works), generous storage.
Cloudflare R2	Audio file storage / CDN	Zero egress. Critical for music delivery.
Anthropic API	AI brain	Tool use is the differentiator.
Stripe Connect	Payment processing	Industry standard for marketplaces.
Resend	Transactional email	Simple, dev-friendly, cheap.
Sentry	Error tracking	Catches bugs before users tweet.
PostHog	Product analytics	Funnel tracking + feature flags.
Cloudflare (free)	DNS + DDoS + WAF	Free belt-and-braces.

Don't do these (common mistakes)

- **✗ AWS from scratch** — you'll spend 3 months on devops before shipping a feature.
- **✗ Self-host Postgres at small scale** — backups + monitoring is more work than it saves.
- **✗ Build your own DDEX delivery in v1** — 6 months of engineering to replicate what RouteNote/SonoSuite has.
- **✗ Use Firebase** — vendor lock-in plus painful migration if you outgrow it.
- **✗ Skip Cloudflare** — even at small scale, free DDoS protection is worth setting up.
- **✗ Run your own email server** — deliverability is a full-time job.

"How much for AI to maintain most of this site?"

Here's the answer to the specific question.

What "AI maintaining the site" means concretely

If you mean "can Claude Code (or similar) handle ongoing engineering tasks like adding features, fixing bugs, updating deps?" — yes, with these caveats:

Things AI can do well today:

- Add new pages / forms / features (90%+ correctness)
- Update libraries to new versions
- Write tests
- Fix UI bugs visible in screenshots
- Generate marketing copy / docs / lesson content
- Refactor code with clear specifications

Things AI cannot do alone:

- Negotiate the RouteNote partnership
- Sign a SOC2 contract
- Decide product direction
- Talk to a real artist about why a feature is broken
- Approve a database migration that drops a column

Cost projection for AI-assisted maintenance

Assuming Claude Code as the primary tool with a human (you or a contractor) reviewing/approving:

Activity	Hours/week	AI cost
Bug fixes + small features	5 hours	25–50/week(100–200/mo)
Larger features (1 new page, deeper integration)	5 hours	50–100/week(200–400/mo)
Dependency updates + housekeeping	1 hour	5/week(20/mo)
Code reviews + planning	2 hours	20/week(80/mo)
Monthly AI cost for maintenance	~13 hours	~\$400–700/mo

Plus a part-time human reviewer (you or contractor) at ~10 hours/week = \$1,500–4,000/mo depending on rate.

Total for AI-assisted maintenance: ~\$2K–5K/mo through Phase 2. At Phase 3+ you need real engineers, not AI.

The honest bound

AI cannot maintain a SaaS that takes payments, processes royalties, and serves DSPs without **at least one human responsible**. Auditors won't accept it. Stripe won't accept it. Spotify won't accept it. Treat AI as a 10× productivity multiplier on a human, not a replacement.

Recommended decision framework

When your boss asks "should we go ahead?" — here's the math:

Break-even analysis (Phase 2 — 500 artists):

- Monthly cost: ~\$1,200
- Per-artist revenue at 15/mo :15
- Break-even at 80 paying artists (15 × 80 =1,200)

Conservative funding ask (12 months runway through Phase 2):

- Monthly burn (no salaries): \$1,200
- One-time setup: \$2,000
- 12 months × 1,200+2,000 + ~30Kengcontractbudget = ** 50K**

With one part-time engineer (you + AI + 1 contractor at \$5K/mo):

- 12 months $\times$ 6,200+2K = **~\$76K**

With one full engineer (\$15K/mo all-in):

- 12 months $\times$ 16K+2K = **~\$194K**
-

Sources

Pricing data current as of April 2026. Always reconfirm before signing:

- [Vercel Pricing](#)
- [Supabase Pricing](#)
- [Stripe Connect Pricing](#)
- [RouteNote Pricing](#)
- [Anthropic API Pricing](#)
- [Cloudflare Stream Pricing](#)
- [US ISRC Agency](#)
- [GS1 US Pricing](#)

For the engineering breakdown (which API, what env var, where to sign up), see [CONNECTIONS_NEEDED.md](#).

For the implementation phases, see [IMPLEMENTATION_PLAN.md](#).